

Exam Solutions

Ahmed Alhasan 'ahmal787'

2020-06-04

Q1

- (a)
- (b)

```
x = 33

logPostTheta <- function(x, theta, a, b){
  sum(dunif(theta,min=a,max = b, log = TRUE) + dbinom(x,size = 50, prob = theta, log = TRUE))
}

a = 0
b = 1

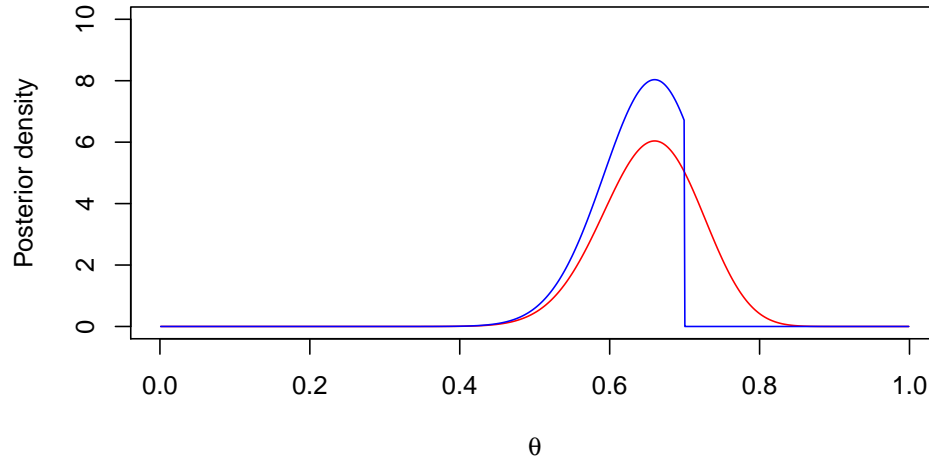
thetaGrid <- seq(0.001, 0.999, length = 1000)
logPostEvals_1 <- rep(0, 1000)
i = 0
for (theta in thetaGrid){
  i = i + 1
  logPostEvals_1[i] <- logPostTheta(x, theta, a, b)
}

a = 0.3
b = 0.7
logPostEvals_2 <- rep(0, 1000)
i = 0
for (theta in thetaGrid){
  i = i + 1
  logPostEvals_2[i] <- logPostTheta(x, theta, a, b)
}

binWidth = thetaGrid[2]-thetaGrid[1]
logPostEvals <- data.frame(Prior_1 = exp(logPostEvals_1)/(sum(exp(logPostEvals_1))*binWidth),
                          Prior_2 = exp(logPostEvals_2)/(sum(exp(logPostEvals_2))*binWidth))

plot(thetaGrid, logPostEvals$Prior_1, type = "l",
     ylab = 'Posterior density',
     xlab = expression(theta),
```

```
col = "red", ylim = c(0,10))
lines(thetaGrid, logPostEvals$Prior_2, type = "l",
      ylab = 'Posterior density',
      xlab = expression(theta),
      col = "blue")
```



- (c)

```
cat("Prior_1: " , mean(logPostEvals$Prior_1<.5))
```

```
## Prior_1: 0.708
```

```
cat("Prior_2: " , mean(logPostEvals$Prior_2<.5))
```

```
## Prior_2: 0.795
```

Q2

- (a)

```
setwd("C:/Machine_Learning/1_Workshop/9_Bayesian_Learning/3_Exams/10_2020_Jun_04")
load(file = 'titanic.RData')

if("mvtnorm" %in% rownames(installed.packages()) == FALSE) {install.packages("mvtnorm")}
if("msm" %in% rownames(installed.packages()) == FALSE) {install.packages("msm")}
library(mvtnorm) # For multivariate normal
library(msm) # For truncated normal

BayesProbReg <- function(y, X, mu_0, tau, nIter){
  # Gibbs sampling in probit regression using data augmentation:
```

```

#
# beta | tau ~ N(mu_0, tau^-2*I)
#
# INPUTS:
# y - n-by-1 vector with response data observations
# X - n-by-nCovs matrix with covariates, first column should be ones if you want an intercept.
# mu_0 - prior mean for beta
# tau - prior standard deviation for beta
# nIter - Number of samples from the posterior (iterations)
#
# OUTPUTS:
# betaSample - Posterior samples of beta.          nIter-by-nCovs matrix

# Prior
nPara <- dim(X)[2]
priorCov <- tau^2*diag(nPara)
priorPrec <- solve(priorCov)

# Compute posterior hyperparameters
n = length(y) # Number of observations
n1 = sum(y)
n0 = n - n1
nCovs = dim(X)[2] # Number of covariates
XX = t(X)%*%X

# The actual sampling
betaSample = matrix(NA, nIter, nCovs)
u <- matrix(NA, n, 1)
beta <- solve(XX, crossprod(X, y)) # OLS estimate as initial value
for (i in 1:nIter){

  xBeta <- X%*%beta

  # Draw u | beta
  u[y == 0] <- rtnorm(n = n0, mean = xBeta[y==0], sd = 1, lower = -Inf, upper = 0)
  u[y == 1] <- rtnorm(n = n1, mean = xBeta[y==1], sd = 1, lower = 0, upper = Inf)

  # Draw beta | u
  betaHat <- solve(XX, t(X)%*%u)
  postPrec <- XX + priorPrec
  postCov <- solve(postPrec)
  betaMean <- solve(postPrec, XX%*%betaHat + priorPrec%*%mu_0)
  beta <- t(rmvnorm(n = 1, mean = betaMean, sigma = postCov))
  betaSample[i,] <- t(beta)

}

return(betaSample=betaSample)
}

y <- titanic[,1]

```

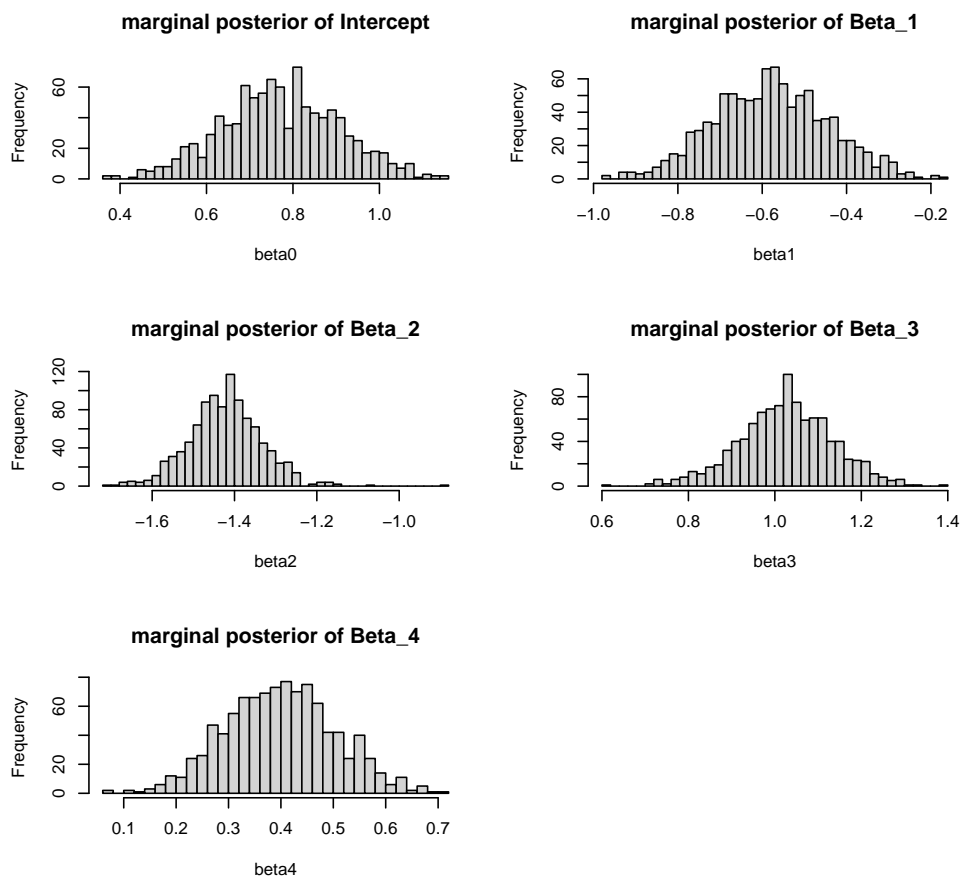
```

X <- as.matrix(titanic[, -1])
mu_0 <- rep(0, dim(X)[2])
tau <- 50
nIter <- 1000

betaSample <- BayesProbReg(y, X, mu_0, tau, nIter)

par(mfrow = c(3,2))
hist(betaSample[,1], breaks = 40, main = "marginal posterior of Intercept", xlab = expression(beta0))
hist(betaSample[,2], breaks = 40, main = "marginal posterior of Beta_1", xlab = expression(beta1))
hist(betaSample[,3], breaks = 40, main = "marginal posterior of Beta_2", xlab = expression(beta2))
hist(betaSample[,4], breaks = 40, main = "marginal posterior of Beta_3", xlab = expression(beta3))
hist(betaSample[,5], breaks = 40, main = "marginal posterior of Beta_4", xlab = expression(beta4))

```



- (b)
- (c)

```

cat("Probability of Beta2 + Beta5>0: ", mean(betaSample[,2] + betaSample[,5] > 0))

```

```
## Probability of Beta2 + Beta5>0: 0.14
```

- (d)

Q3

- (a)

```
x <- c(log(20), log(20), log(50), log(40))
y <- c(5, 3, 17, 8)
logPostBeta <- function(beta, x, y){
  PostBeta <- vector(length = 4)
  for(i in 1:length(x)){
    bPrior <- dnorm(x[i], 1, 0.1, log = TRUE)
    logLike <- sum(dpois(y[i], lambda = exp(x*beta), log = TRUE))
    PostBeta[i] <- bPrior + logLike
  }
  return(PostBeta)
}

initVal <- as.vector(rep(0,length(x)))
OptimResults <- optim(initVal,logPostBeta, gr=NULL, x, y,method=c("BFGS"),
                      control=list(fnscale=-1),hessian=TRUE)
```

- (b)

Q4

- (c)

```
prob_fm <- dnorm(10, 14, 2) * dnorm(250, 300, 50) * 5/22
prob_fm
```

```
## [1] 2.969144e-05
```